

# CAI 4104: Machine Learning Engineering

Project Report: **Object Classification**

Khang Tran  
trankhang@ufl.edu

Nicolas Slenko  
nslenko@ufl.edu

Colin Wishart  
c.wishart@ufl.edu

Christopher Dowdy  
cdowdy@ufl.edu

Mauro Chavez Vega  
maurochavezvega@ufl.edu

April 25, 2025

## 1 Introduction

Given a dataset containing 128 x 128 RGB PNG images containing one of 12 different objects: backpack, book, calculator, chair, clock, desk, keychain, laptop, paper, pen, phone, water bottle. The goal of this project is to correctly classify images into one of the 12 classes mentioned above. The dataset is split into 70% training, 15% validation and 15% testing. We will use a convolutional neural network (CNN) to classify the images. The CNN will consist of several convolutional layers, followed by max pooling layers, and finally fully connected layers. We will use the Adam optimizer and categorical cross-entropy loss function to train the model. The model will be evaluated using accuracy and per class accuracy.




Figure 1: Distribution of images across the 12 object classes in the dataset. The dataset is relatively balanced, with each class containing between 360-420 images.

## 2 Approach

### 2.1 Data Preprocessing

Our data preprocessing pipeline consists of several key steps to prepare the images for training:

- **Image Resizing:** Although it is assumed that all the images are 128x128 pixels, we still resized all the images to 128x128 pixels to ensure uniformity. and that the expected batch size match the input size of the model.

- **Normalization:** The input were first applied with transform.ToTensor which would have performed scaling on the pixel values from the range of [0, 255] to [0, 1]. All images were then normalized using mean = [0.5, 0.5, 0.5] and standard deviation = [0.5, 0.5, 0.5] for each color channel to center the data around 0 and remap it to a range of [-1, 1] and ensure consistent scale.
- **Data Splitting:** We divided the dataset into three subsets:
  - Training set: 70% of the data
  - Validation set: 15% of the data
  - Test set: 15% of the data

We decided on that this data splitting ratio since that is what was recommended for small dataset (under 100k examples) such as the dataset we are using to train this project model.

## 2.2 Model Architecture

Since the problem at hand is clearly a multi-class classification problem, we decided to use a convolutional neural network (CNN) architecture which generally performs well on computer vision tasks or task related to image data. We implemented a custom CNN architecture with the following layers:

- **Data Augmentation:** Since the given dataset only consists of 4757 images, which is relatively small, we applied data augmentation to just the training set to artificially increase training diversity. Furthermore, this also act as a regularization technique to prevent overfitting which we later encountered during training. The data augmentation techniques we used include:
  - Random horizontal flips to 50 percent of the images
  - Random vertical flips to 50 percent of the images
  - Random rotations (up to 15 degrees) to 50 percent of the images

These techniques should help the model generalize better to unseen data as the objects in real life might appear in different orientations.

- **Convolutional Layers:** Five convolutional blocks, each containing:
  - A convolutional layer with increasing filter depths (starting from base filters and increasing)
  - Batch normalization for further improved training stability
  - ReLU activation function
  - Max pooling layer
  - Dropout for regularization

We decided to go with a stretch pants like approach where we made our layers wider and then use dropout in between them.

- **Fully Connected Layers:** Two fully connected layers:
  - A flatten layer that reshape the 512x4x4 feature map into a 8192-dimensional vector
  - A hidden layer with dropout regularization of 0.5 to prevent overfitting
  - An output layer with 12 neurons (one for each class)
- **Loss Function:** Cross-entropy loss since it is a classification problem with label smoothing of 0.1, which helps prevent the model from becoming overconfident and reduce overfitting
- **Optimizer:** Adam optimizer with a learning rate of 0.0005 and weight decay (L2 regularization) of 0.0005

This architecture was designed to extracting higher-level features through the deep convolutional layers for good performance while still preventing overfitting through the use of regularization techniques.

## 2.3 Training Strategy

- **Learning Rate Scheduling:** We implemented a ReduceLROnPlateau scheduler that reduces the learning rate by a factor of 0.5 when validation accuracy plateaus for 4 consecutive epochs
- **Early Stopping:** Training terminates if validation loss doesn't improve for 10 consecutive epochs, preventing overfitting
- **Model Checkpointing:** We saved both an early-stopped model and the final model after training completion for faster evaluation and testing

## 2.4 Training Process Visualization

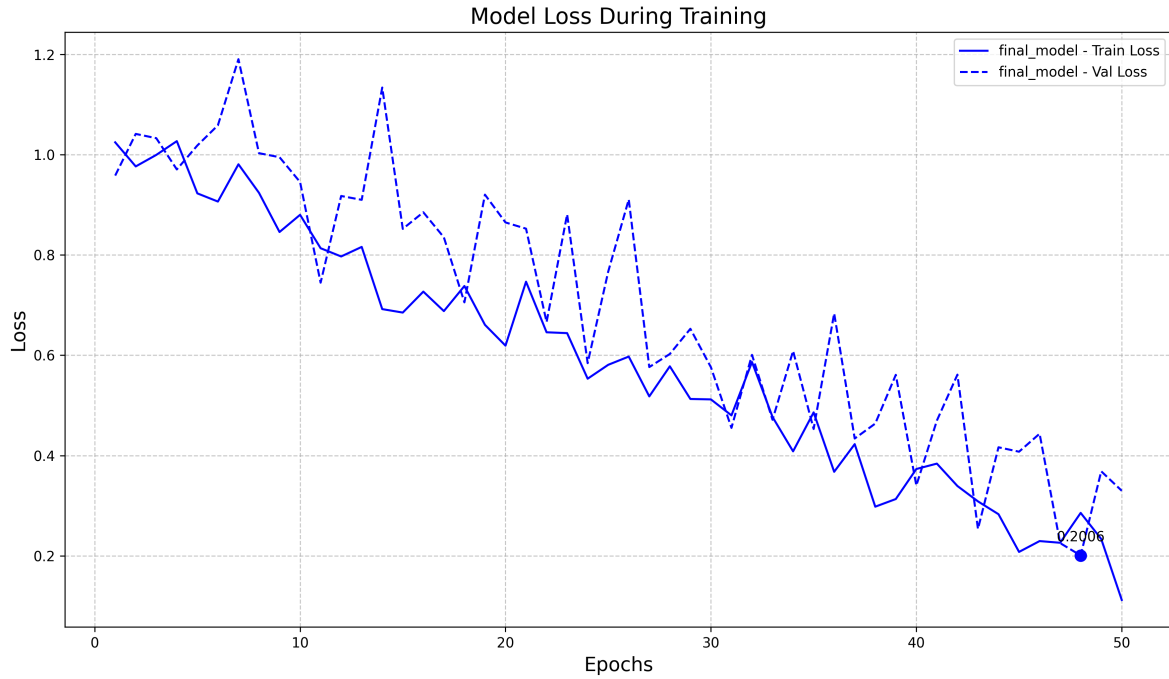


Figure 2: Model loss during training. The training loss (solid line) and validation loss (dashed line) both decrease over time, indicating successful learning.




Figure 3: Model accuracy during training. The training accuracy (solid line) and validation accuracy (dashed line) both increase over time, with final validation accuracy reaching approximately 88.85%. Note that this validation accuracy differs from the 78.43% test accuracy reported in Table 2, as they are evaluated on different datasets.

### 3 Evaluation Methodology

- **Metrics:** Since this is a multi-class classification problem, we used the following metrics to evaluate our model:
  - Accuracy: To measure the overall performance of the model
  - Per-class accuracy: The accuracy for each class to evaluate the model’s performance on each class
- **Baselines:** We compared our model’s performance by evaluating the metrics mentioned above against a random guessing model as a baseline.
- **Data Split:** As mentioned in our preprocessing section, we used a 70%-15%-15% train-val-test split on our dataset.
- **Hyperparameter Tuning:** The hyperparameters that were tuned in our model include:
  - Different learning rates (0.0001, 0.0005, 0.001)
  - Different batch sizes (32, 64, 128) with 64 being the most optimal for learning speed and computation time
  - Different dropout rates in both convolutional (0.1) and fully connected (0.5) layers
  - Different weight decay values for regularization (0.0001, 0.0005, 0.001)
  - Different numbers of convolutional blocks (3, 5, 7) with 5 being the most optimal for learning speed and computation time
- **Model Selection:** The final model was selected based on the best accuracy and per class accuracy on the test set. Additionally, we also considered the model’s performance on the validation set to ensure that overfitting was at an acceptable level.

## 4 Results

### 4.1 Test-Set Composition

The test set contains a total of 714 images distributed among the 12 classes as follows:

Table 1: Test-set image counts by class

Class	Image Count
Backpack	68
Book	53
Calculator	65
Chair	60
Clock	57
Desk	55
Keychain	59
Laptop	60
Paper	57
Pen	47
Phone	73
Water Bottle	60
<b>Total</b>	<b>714</b>

## 4.2 Accuracy Comparison

Table 2: Per-class and overall performance: CNN vs. random baseline

Class	CNN Accuracy (%)	Random Baseline (%)	Improvement (%)
Backpack	76.47	7.35	69.12
Book	77.36	7.55	69.81
Calculator	87.69	9.23	78.46
Chair	65.00	6.67	58.33
Clock	87.72	3.51	84.21
Desk	60.00	7.27	52.73
Keychain	84.75	6.78	77.97
Laptop	80.00	10.00	70.00
Paper	87.72	12.28	75.44
Pen	89.36	6.38	82.98
Phone	68.49	6.85	61.64
Water Bottle	80.00	6.67	73.33
<b>Overall</b>	<b>78.43</b>	<b>7.56</b>	<b>70.87</b>

## 4.3 Results Analysis

Table 1 already shows that our CNN model achieves an overall test accuracy of 78.43% versus only 7.56% for the random guessing baseline an improvement of 70.87 percentage points. A more detailed analysis of how the models performance on different classes differs is as the following:

## 4.4 Analysis of Per-Class Performance

**Highest Gains** The largest improvements over random chance occur on:

- **Clock:** +84.21 pts (87.72% vs. 3.51%)
- **Pen:** +82.98 pts (89.36% vs. 6.38%)
- **Calculator:** +78.46 pts (87.69% vs. 9.23%)

**Lower Gains** Smaller improvements are seen on:

- **Desk:** +52.73 pts (60.00% vs. 7.27%)
- **Chair:** +58.33 pts (65.00% vs. 6.67%)
- **Phone:** +61.64 pts (68.49% vs. 6.85%)

**Effect of Test Set Distribution** Analyzing the per-class performance in relation to the number of images in each class shows that there is no direct correlation between class size and classification accuracy.

- *Phone*, the most common class (73 images), achieved a modest accuracy of 68.49%, which is lower than average.
- *Pen*, the least represented class (47 images), achieved the highest accuracy of 89.36%.

This suggests that class frequency on its own does not determine model performance. It can potentially be explained that having clearer, more distinctive visual features (e.g., pens, clocks, calculators) is what really influence CNN accuracy to be higher instead of just the number of examples in each class.

**Consistently Strong Classes** Several mid-sized classes performed particularly well:

- **Clock (57 images)** and **Calculator (65 images)** both achieved nearly 88% accuracy.
- **Paper (57 images)** and **Keychain (59 images)** also scored above 84%.

These results further confirms the idea that the CNN benefits from learning distinct shapes and textures, as many of these objects have well-defined and distinct appearance.

**Weaker Classes and Ambiguity**

- **Desk (55 images):** 60.00%
- **Chair (60 images):** 65.00%

Classes with lower accuracy like *Desk* and *Chair* suggest that visual ambiguity may hinder classification, even when test counts are moderate. The given images of desks or chairs may have appeared in a wider range of angles, lighting conditions compared to other classes. Furthermore, the many different designs and appearances that they might come in could have also made it harder for the model to generalize.

## 4.5 Summary

Regardless of a few underperforming classes, our model has significantly outperformed the random guessing baseline by 70.87 percentage points overall, reaching a 78.43% accuracy, with every class achieving more than a 50-point gain. This demonstrates that our CNN model is effective at classifying images into the 12 object classes.

## 5 Conclusions

Write here.

## References